



Can you recognise the music playing?

6COSC020W: APPLIED AI

WEEK 8: ARTIFICIAL NEURAL NETWORKS (ANNs)

DR. ESTER BONMATI


(e.bonmaticoll@westminster.ac.uk)

NOTES

- We will use PollEverywhere multiple times during the lecture: polleverywhere.com/esterbonmati
- It may be useful to have a notebook and pen/pencil (or your preferred method) to take some notes.

Before session: What do you think machine learning is?

FROM LAST SESSION...



LEARNING OUTCOMES OF THIS SESSION

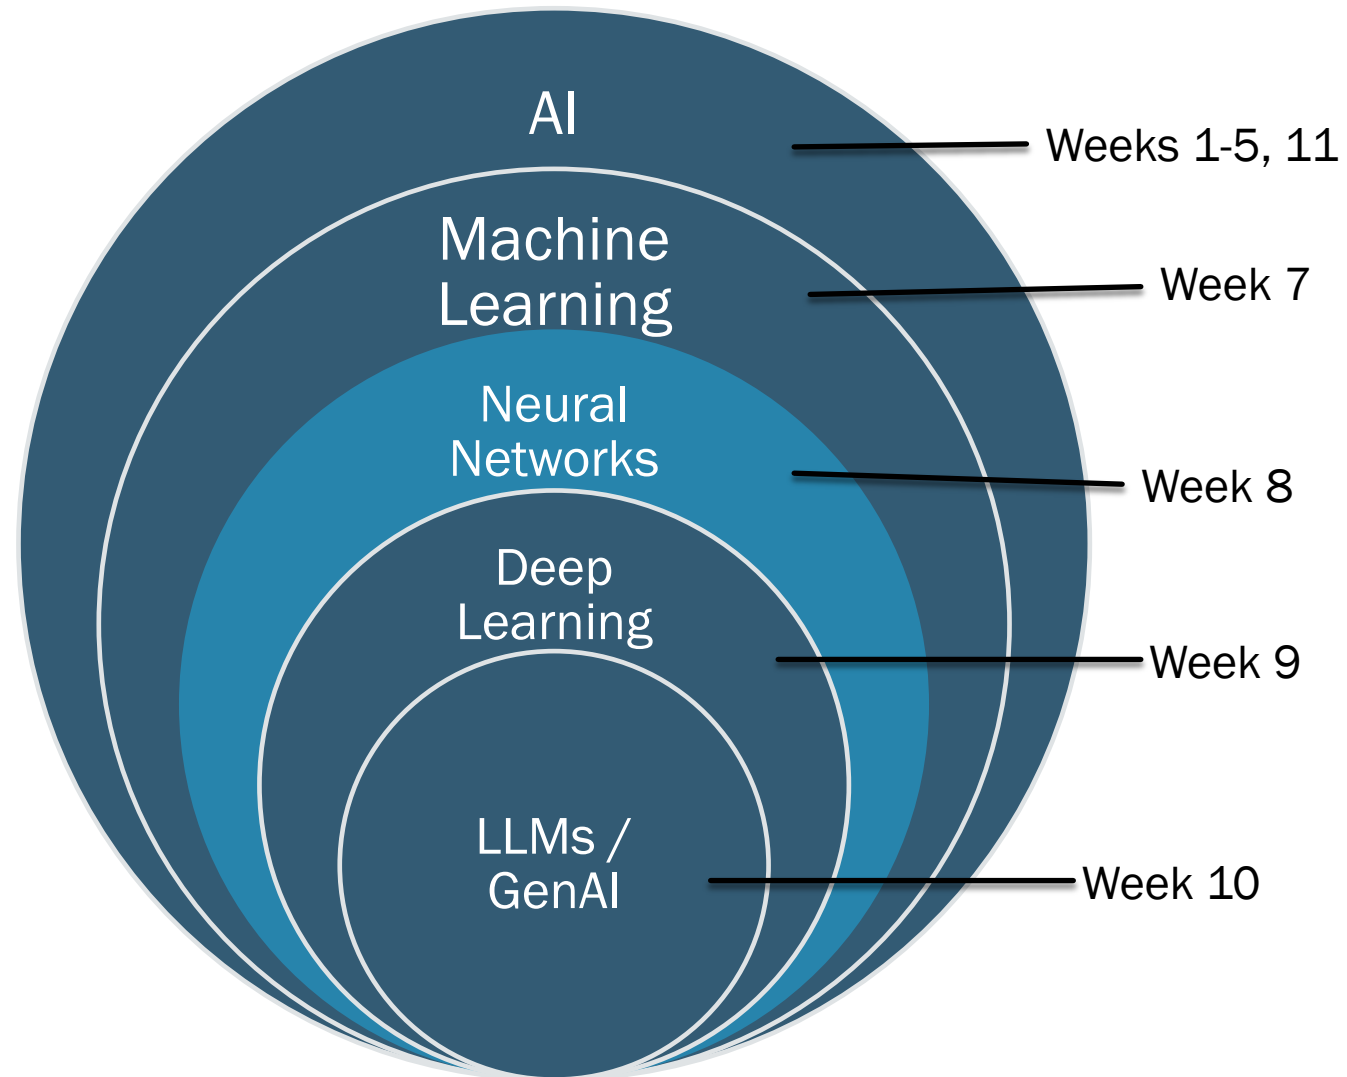
Perceptron

- Explain the perceptron
- Inputs and outputs
- Weights
- Forward propagation
- Explain what is an activation functions
- Be able to calculate the output of a perceptron

Neural Networks

- Be able to list the name of the different layers
- Be able to describe how neural networks work in your own words
- Explain problems with data
- Explain what are hyperparameters
- Be able to evaluate performance and calculate accuracy

ARTIFICIAL INTELLIGENCE: MACHINE LEARNING



Before: what is an artificial neural network?

pollev.com/esterbonmati

Nobody has responded yet.

Hang tight! Responses are coming in.

Do you know any examples of applications using a neural network?

pollev.com/esterbonmati

Nobody has responded yet.


Hang tight! Responses are coming in.

ARTIFICIAL NEURAL NETWORKS (ANNs)

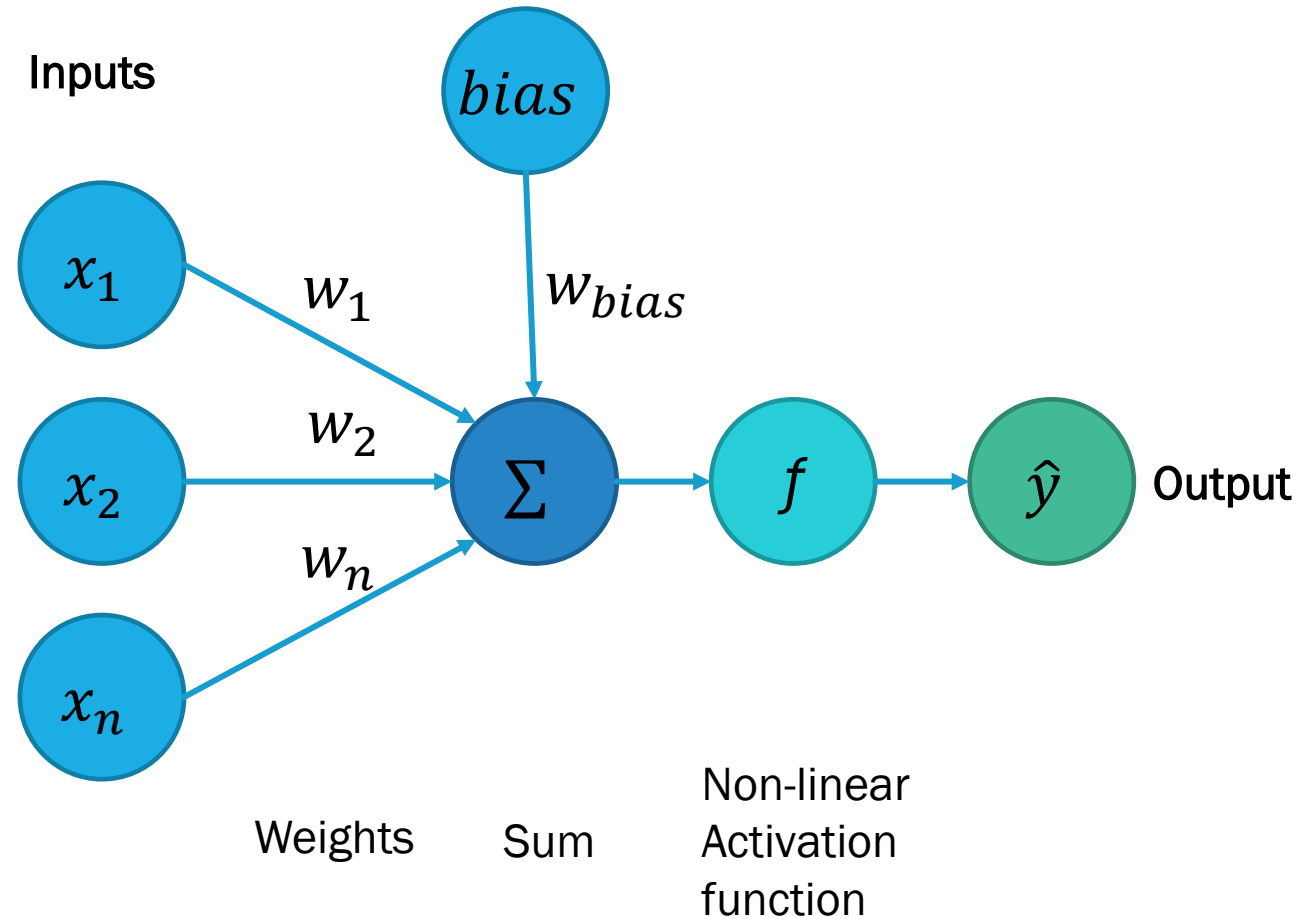
Why are they important?
ANN can **help computers**
make intelligent
decisions with limited
human assistance.
They can learn and
model the **relationships**
between input and
output data that are
nonlinear and complex.



NEURON




THE PERCEPTRON



THE PERCEPTRON

- Building block of neural networks (single neuron)
- Mimics the neuron in the human brain
- Forward propagation (input → neuron → output)



$$y = f \left(bias \, w_{bias} + \sum_{i=1}^n x_i \, w_i \right)$$

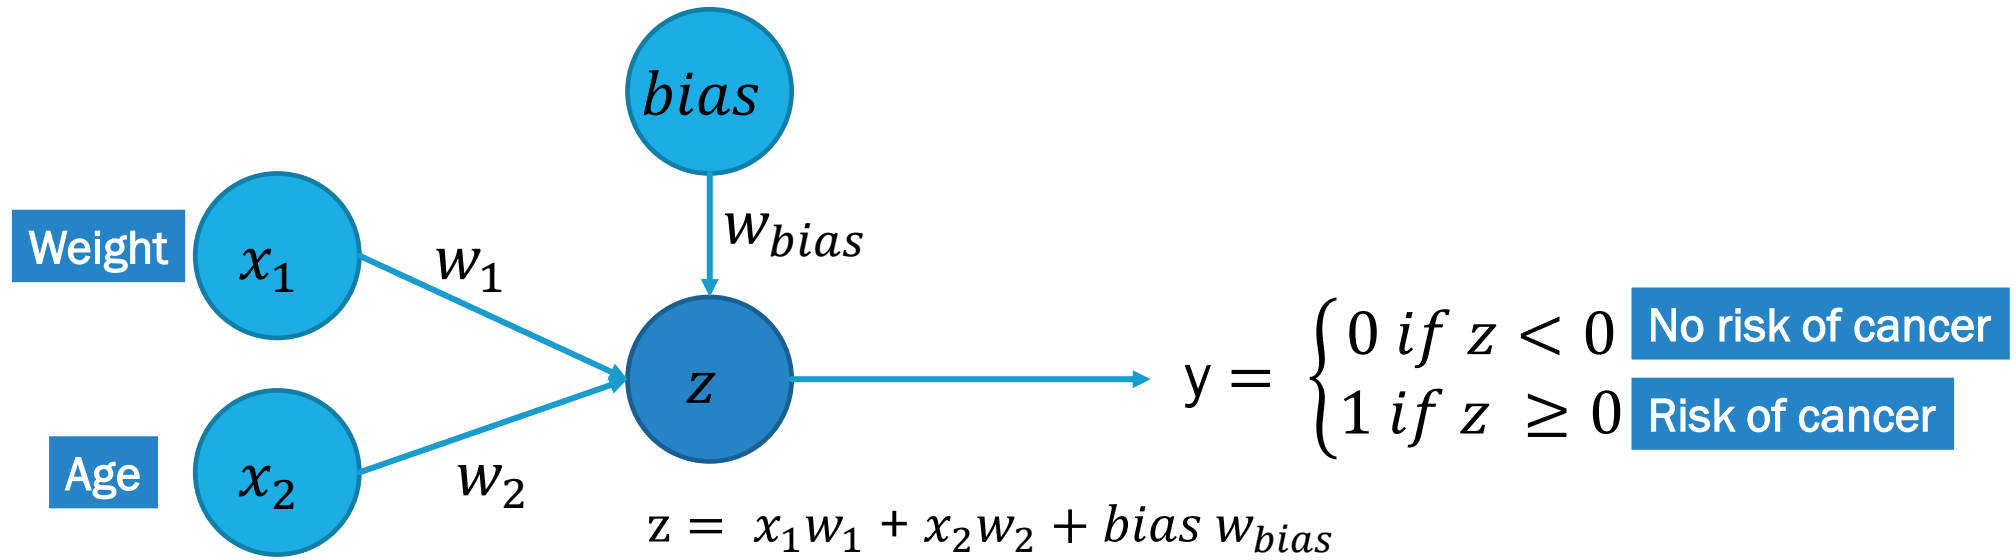
Diagram illustrating the mathematical formula for the perceptron output y . The formula is annotated with arrows indicating the components: y is the Output; f is the Non-linear function; $bias$ is the Bias; w_{bias} is the weight for the bias; $\sum_{i=1}^n$ is the Sum; x_i are the Inputs; and w_i are the Weights.

$$z = x_1 \, w_1 + x_2 \, w_2 + \dots + x_n \, w_n + bias \, w_{bias}$$

$$y = f(z)$$

$$\text{Example: } y = \begin{cases} 0 & \text{if } z < 0 \\ 1 & \text{if } z \geq 0 \end{cases}$$

PERCEPTRON: FORWARD PROPAGATION EXAMPLE




Do you know another example of a perceptron?

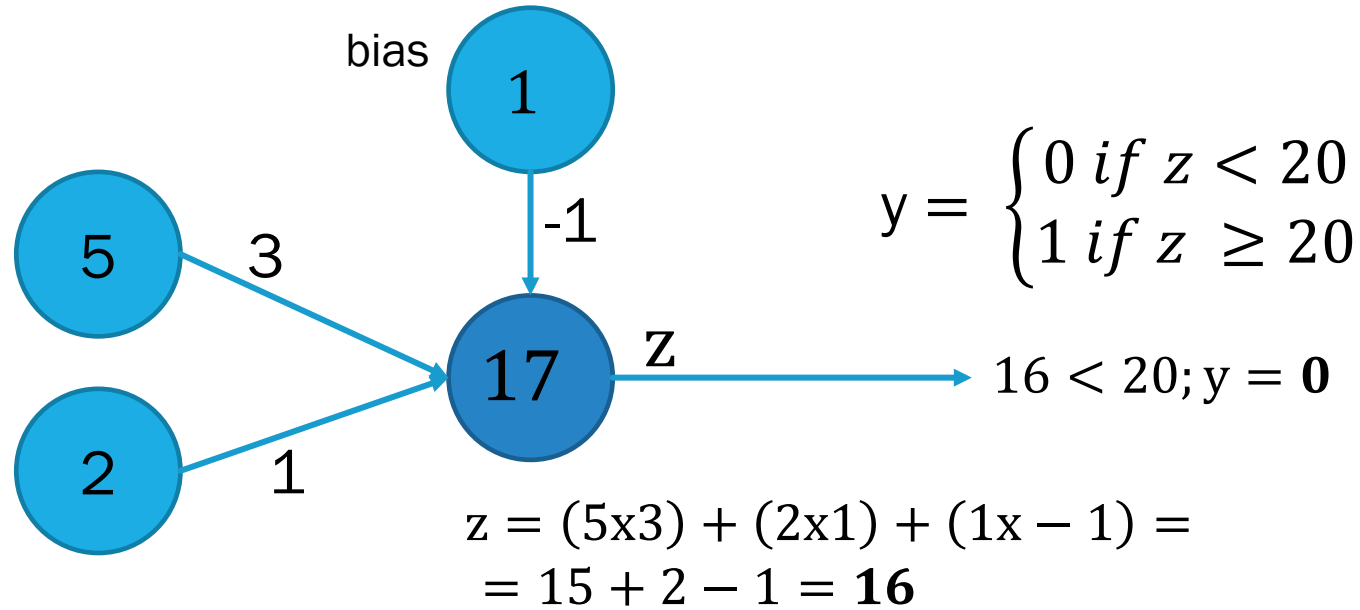
- Logic operations (except XOR)
- Linear separation
- Predict who is going to pass this module based on the number of lectures and tutorials attended



PERCEPTRON: FORWARD PROPAGATION EXAMPLE



PERCEPTRON: FEEDBACK



pollev.com/esterbonmati

When poll is
active respond
at

PollEv.com
/esterbonmati




Result perceptron 1

Nobody has responded yet.

Hang tight! Responses are coming in.

Compare answers with the person sitting next to you

PERCEPTRON: FEEDBACK



Compare answers with the person sitting next to you

pollev.com/esterbonmati
When poll is active respond at

PollEv.com
/esterbonmati




Results perceptron 2

Nobody has responded yet.

Hang tight! Responses are coming in.

PERCEPTRON: HOW DO WE FIND THE WEIGHTS?

Problem: we want to train a perceptron to classify points in 2 classes.



How do we find the weights to classify the points?

$$z = x_1 w_1 + x_2 w_2 + x_0 w_0$$

PERCEPTRON: INPUT DEFINITION



Problem: We have 2 groups of samples, each of them defined by 2 values, and we want to use a perceptron to classify them into their group.

```
x1_group1 = [2, 5, 4, 2.5, 3.5, 4.5]
```

```
x2_group1 = [1, 2, 3, 2, 1, 4]
```


```
x1_group0 = [0.5, 2.5, 3, 1.5, 1, 2]
```

```
x2_group0 = [3, 5.5, 4.5, 3.5, 4, 5]
```



PERCEPTRON: ACTIVATION FUNCTION

$$y = \begin{cases} -1 & \text{if } z \leq 0 \\ 1 & \text{if } z > 0 \end{cases}$$



If the sum is greater than 0 we return 1, otherwise we return 0


```
In [112]: def activation_function(value):  
          if value > 0: return 1  
          else: return -1
```

PERCEPTRON: WEIGHTS AND BIAS DEFINITION

Some random weights...

```
w1 = 0.75  
w2 = -0.5  
bias = 1  
wbias = 1
```


How good are these weights in grouping the input data?



PERCEPTRON: FORWARD PROPAGATION

We need to find z in order to calculate the output y ...

```
def perceptron(x1, x2, bias):  
    z = x1*w1 + x2*w2 + bias*wbias  
    return z
```



Then we can calculate the output y using the perceptron and the activation function:

```
for i in range(len(x1_group0)):  
    z = perceptron(x1_group0[i], x2_group0[i], bias)  
    group = activation_function(z)  
    print('Point ' + str(i) + ' in group 0 returned ' + str(group))
```

PERCEPTRON: PREDICTION AND RESULTS


We calculate the output for each group, so we know which is the prediction:

```
group = activation_function(z)
print('Point ' + str(i) + ' in group 1 returned ' + str(group))

for i in range(len(x1_group0)):
    z = perceptron(x1_group0[i], x2_group0[i], bias)
    group = activation_function(z)
    print('Point ' + str(i) + ' in group 0 returned ' + str(group))
```

```
Point 0 in group 1 returned 1
Point 1 in group 1 returned 1
Point 2 in group 1 returned 1
Point 3 in group 1 returned 1
Point 4 in group 1 returned 1
Point 5 in group 1 returned 1
Point 0 in group 0 returned -1
Point 1 in group 0 returned 1
Point 2 in group 0 returned 1
Point 3 in group 0 returned 1
Point 4 in group 0 returned -1
Point 5 in group 0 returned -1
```

} Incorrect




$$accuracy = \frac{\text{number correct samples}}{\text{total number of samples}} = \frac{9}{12} = 0.75 \rightarrow 75 \%$$

Can we improve the accuracy if we change the weights?

PERCEPTRON: VISUALISATION

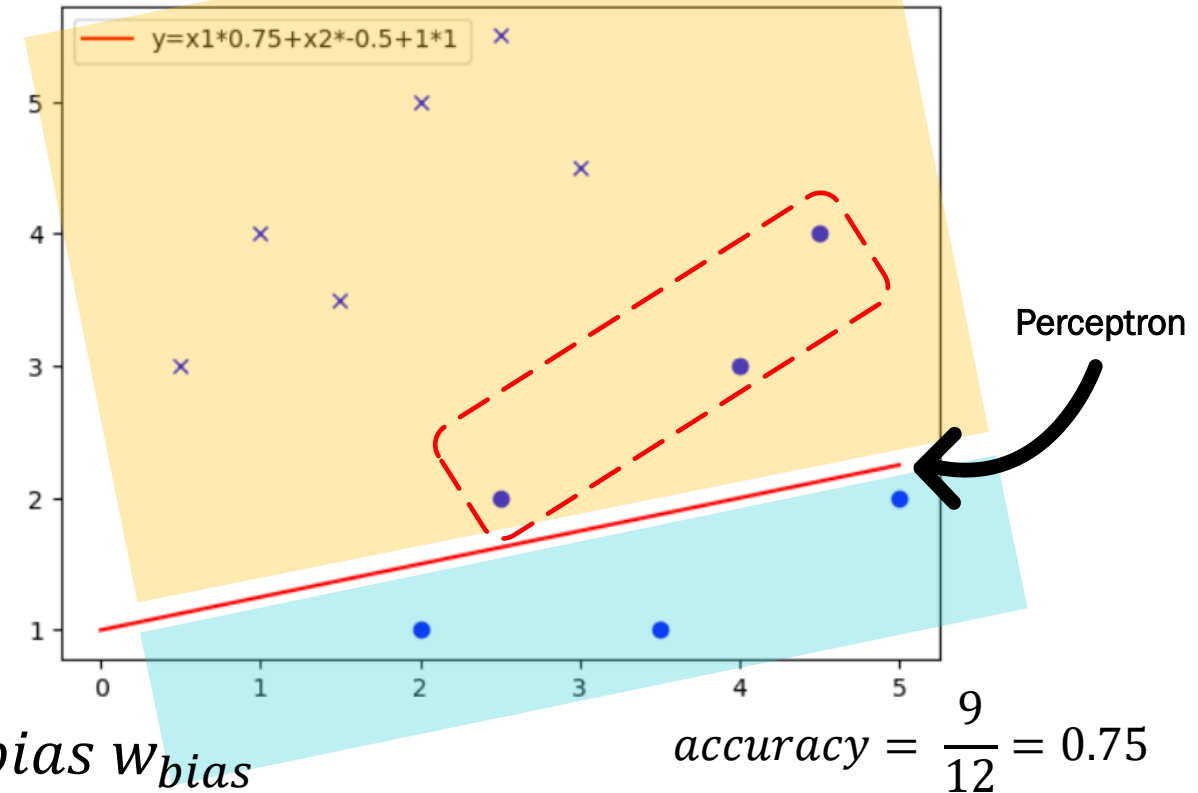
Plot the points

```
In [117]: fig = plt.figure()
ax = plt.axes()
plt.plot(x1_group1, x2_group1, 'ob')
plt.plot(x1_group0, x2_group0, 'xb')
plt.show()
```



```
[7]: x1 = np.linspace(0,5,100)
x2 = np.linspace(0,5,100)
y_1 = w1*x1 + w2*x2 + 1

plt.plot(x1_group1, x2_group1, 'ob')
plt.plot(x1_group0, x2_group0, 'xb')
plt.plot(x1, y_1, '-r', label='y=x1*0.75+x2*-0.5+1*1')
plt.legend(loc='upper left')
plt.show()
```




PERCEPTRON: UPDATE WEIGHTS

To change the prediction output, we need to change the weights

```
In [102]: w1 = 1
w2 = 1
bias = 1
wbias = 1

y_1 = w1*x1 + w2*x2 + bias*wbias

plt.plot(x1_group1, x2_group1, 'ob')
plt.plot(x1_group0, x2_group0, 'xb')
plt.plot(x1, y_1, '-r', label='y=x1*1+x2*1+1*1')
plt.legend(loc='upper left')
plt.show()
```




$$acc = \frac{8}{12} = 0.67$$

```
In [103]: w1 = 3
w2 = -2
bias = 1
wbias = -1

y_1 = x1*w1 + x2*w2 + bias*wbias

plt.plot(x1_group1, x2_group1, 'ob')
plt.plot(x1_group0, x2_group0, 'xb')
plt.plot(x1, y_1, '-r', label='y=x1*3+x2*-2+1*-1')
plt.legend(loc='upper left')
plt.show()
```



$$acc = \frac{12}{12} = 1.0$$

PERCEPTRON: NEW PREDICTION

```
In [104]: for i in range(len(x1_group1)):
          sum = perceptron(x1_group1[i], x2_group1[i], bias)
          group = activation_function(sum)
          print('Point ' + str(i) + ' in group 1 returned ' + str(group))


          for i in range(len(x1_group0)):
              sum = perceptron(x1_group0[i], x2_group0[i], bias)
              group = activation_function(sum)
              print('Point ' + str(i) + ' in group 0 returned ' + str(group))
```

```
Point 0 in group 1 returned 1
Point 1 in group 1 returned 1
Point 2 in group 1 returned 1
Point 3 in group 1 returned 1
Point 4 in group 1 returned 1
Point 5 in group 1 returned 1
Point 0 in group 0 returned -1
Point 1 in group 0 returned -1
Point 2 in group 0 returned -1
Point 3 in group 0 returned -1
Point 4 in group 0 returned -1
Point 5 in group 0 returned -1
```





$$acc = \frac{12}{12} = 1.0 \rightarrow 100 \%$$

NEURAL NETWORKS: ANOTHER EXAMPLE



NEURAL NETWORKS: ANOTHER EXAMPLE

$$y = x_1w_1 + x_2w_2 + \boxed{x_3w_3} + b$$



[For reference only] Equation of a plane = $ax + by + cz + d$






Join by Web Pollev.com/esterbonmati

pollev.com/esterbonmati



How many inputs can a perceptron have?

2



0%

3

SEE MORE






When poll is active respond at [PollEv.com/esterbonmati](https://pollev.com/esterbonmati)

pollev.com/esterbonmati



How many weights has each input in a perceptron?

0



0%

1

SEE MORE



When poll is active respond at [PollEv.com/esterbonmati](https://pollev.com/esterbonmati)

pollev.com/esterbonmati

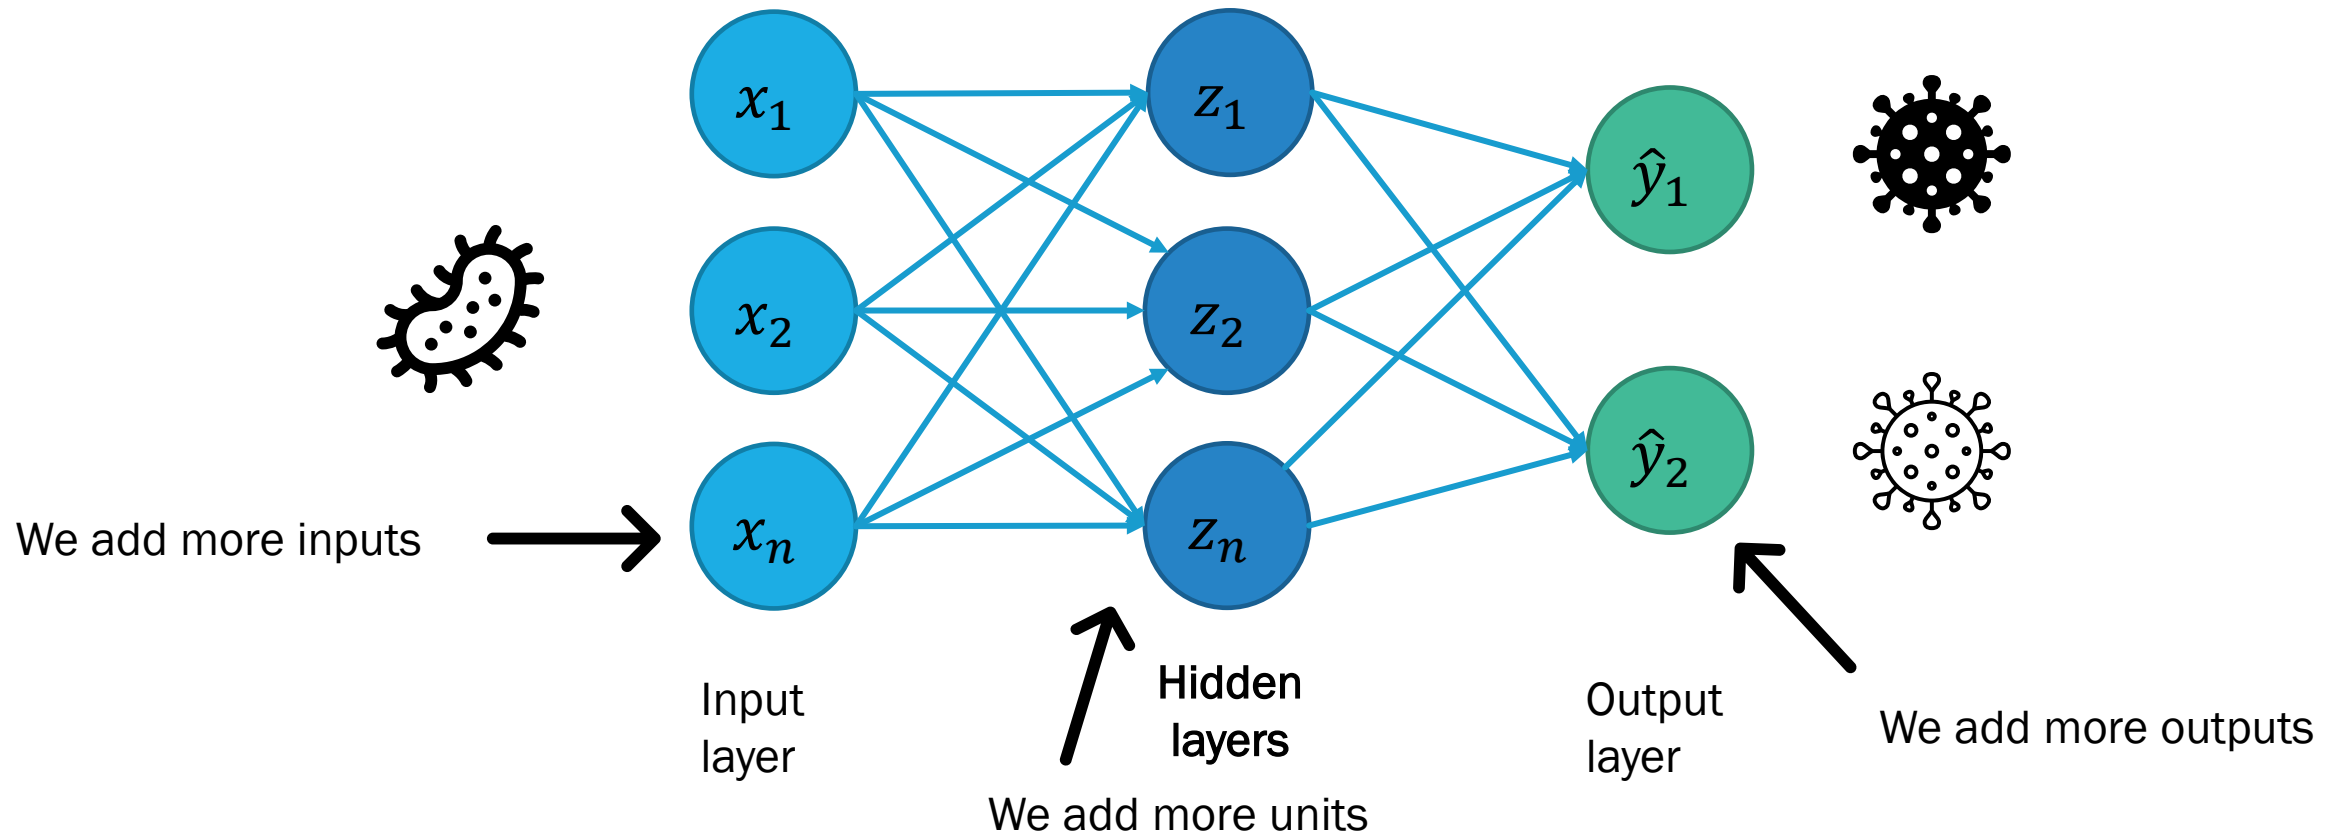


Q&A


Nobody has responded yet.

Hang tight! Responses are coming in.


MULTI-LAYER NEURAL NETWORKS



MULTI-LAYER NEURAL NETWORKS EXAMPLE



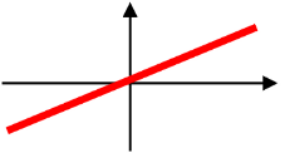
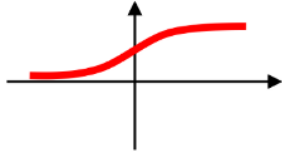
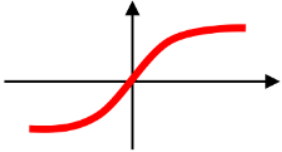
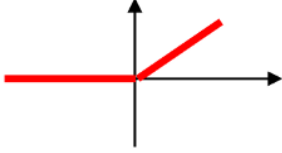
ARTIFICIAL NEURAL NETWORKS (ANN)



Each output has its own weight for each input

$$z = w_0 + \sum_{i=1}^n x_i w_i$$
$$\hat{y} = f \left(w_0 + \sum_{i=1}^n x_i w_i \right)$$

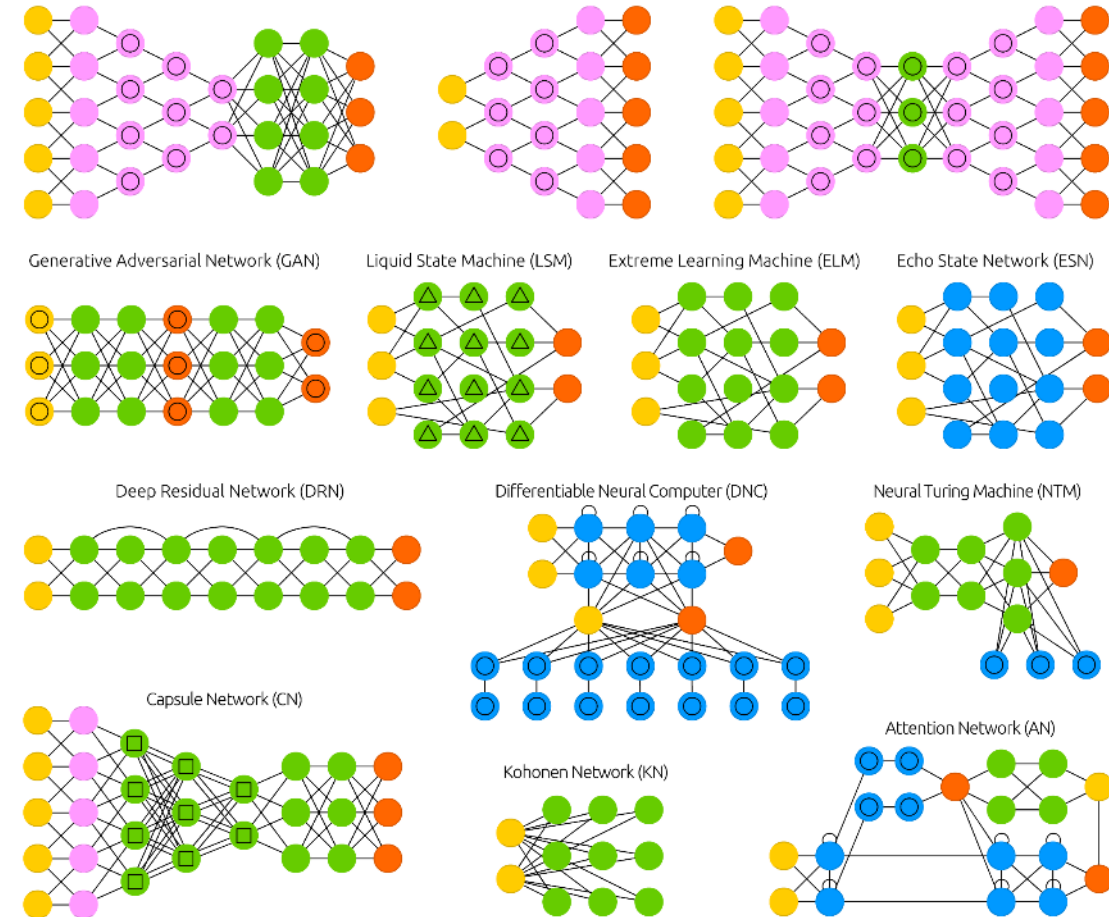
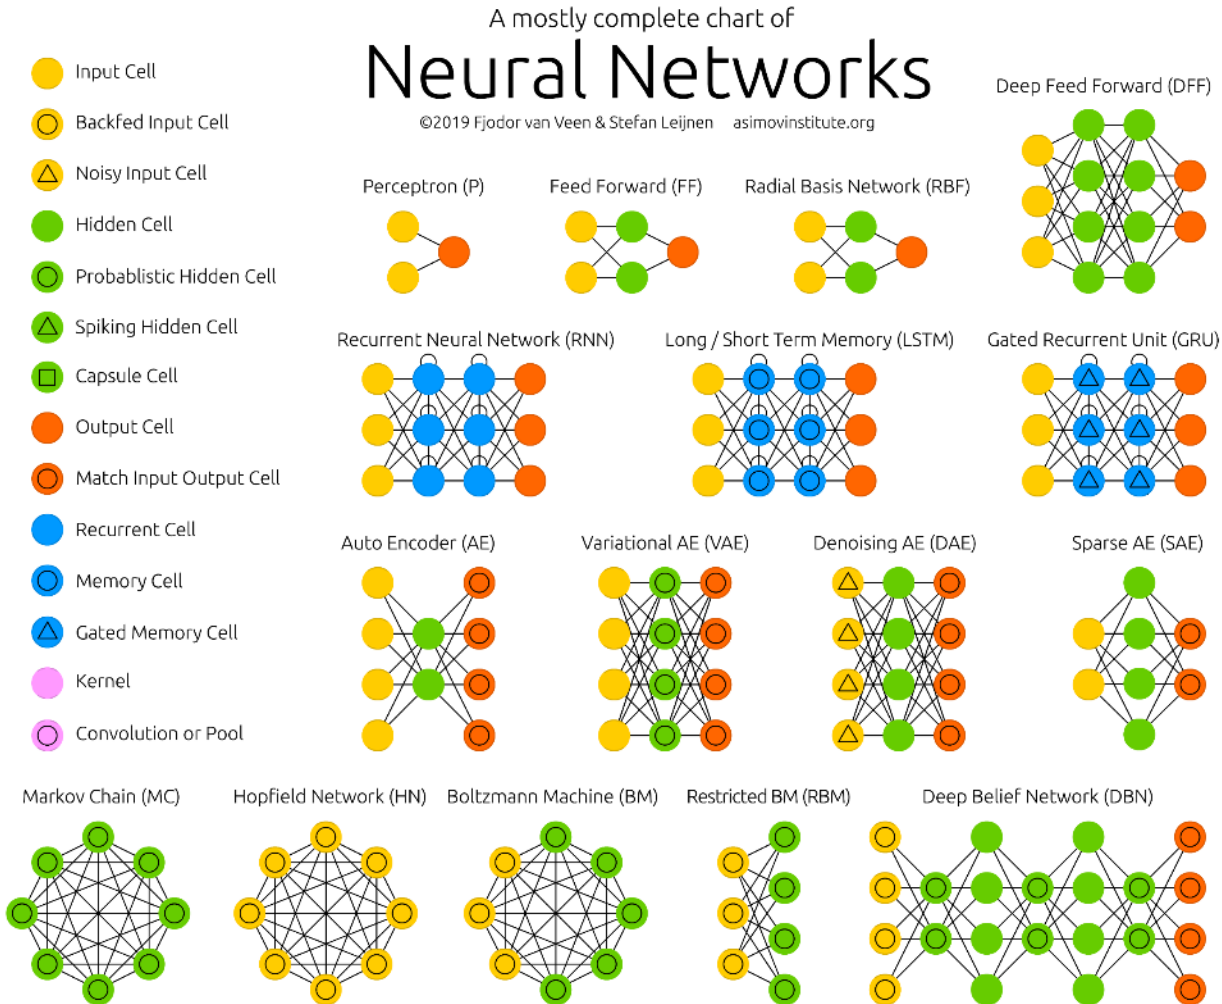
FOR REFERENCE ONLY: ACTIVATION FUNCTIONS

Activation function	Equation	Example	1D Graph
Linear	$\phi(z) = z$	Adaline, linear regression	
Non-linear	Logistic (sigmoid)	Logistic regression, Multi-layer NN	
	Hyperbolic tangent (Tanh)	Multi-layer Neural Networks	
	Rectifier, ReLU (Rectified Linear Unit)	Multi-layer Neural Networks	


Non-linear

Non-linear activation functions allow to learn more complex mappings

EXAMPLES OF NEURAL NETWORKS



WHAT CAN WE DO WITH ANN?



Shopping prediction



Recognition/Classification

Interests prediction



Chatbots

Drug discovery



Stock market prediction

Do you know more examples?



RESEARCH

Forecasting Climatic Trends Using Neural Networks: An Experimental Study Using Global Historical Data

Takeshi Ise^{1,2*} and Yurika Oba¹

¹ Field Science Education and Research Center (FSERC), Kyoto University, Kyoto, Japan, ² Japan Science and Technology Agency (JST), Kawaguchi, Japan

Review Article

Artificial neural networks in the cancer genomics frontier

Andrew Oustimov¹, Vincent Vu²

¹Department of Epidemiology & Biostatistics, College of Public Health, University of South Florida, Tampa, FL 33620, USA; ²Department of Mathematics and Statistics, University of California, Los Angeles, CA, USA

Correspondence to: Andrew Oustimov, MPH. Department of Epidemiology & Biostatistics, College of Public Health, University of South Florida, 13201 Bruce B Downs Blvd, Tampa, FL 33620, USA. Email: aoustimo@mail.usf.edu.

IEEE TRANSACTIONS ON NEURAL NETWORKS AND LEARNING SYSTEMS, VOL. 32, NO. 10, OCTOBER 2021

4291

Attention in Natural Language Processing

Andrea Galassi¹, Marco Lippi¹, and Paolo Torrioni¹

AUTOMATIC LANGUAGE IDENTIFICATION USING DEEP NEURAL NETWORKS

Ignacio Lopez-Moreno¹, Javier Gonzalez-Dominguez^{1,2}, Oldrich Plchot³, David Martinez⁴,
Joaquin Gonzalez-Rodriguez², Pedro Moreno¹

¹Google Inc., New York, USA

²ATVS-Biometric Recognition Group, Universidad Autonoma de Madrid, Spain

³Brno University of Technology, Czech Republic

⁴Aragon Institute for Engineering Research (I3A), University of Zaragoza, Spain

{elnota, jgd}@google.com

Tran et al. *Genome Medicine* (2021) 13:152
<https://doi.org/10.1186/s13073-021-00968-x>

Genome Medicine

REVIEW


Open Access

Deep learning in cancer diagnosis, prognosis and treatment selection

Khoa A. Tran^{1,2}, Olga Kondrashova¹, Andrew Bradley⁴, Elizabeth D. Williams^{2,3}, John V. Pearson¹ and Nicola Waddell^{1*}

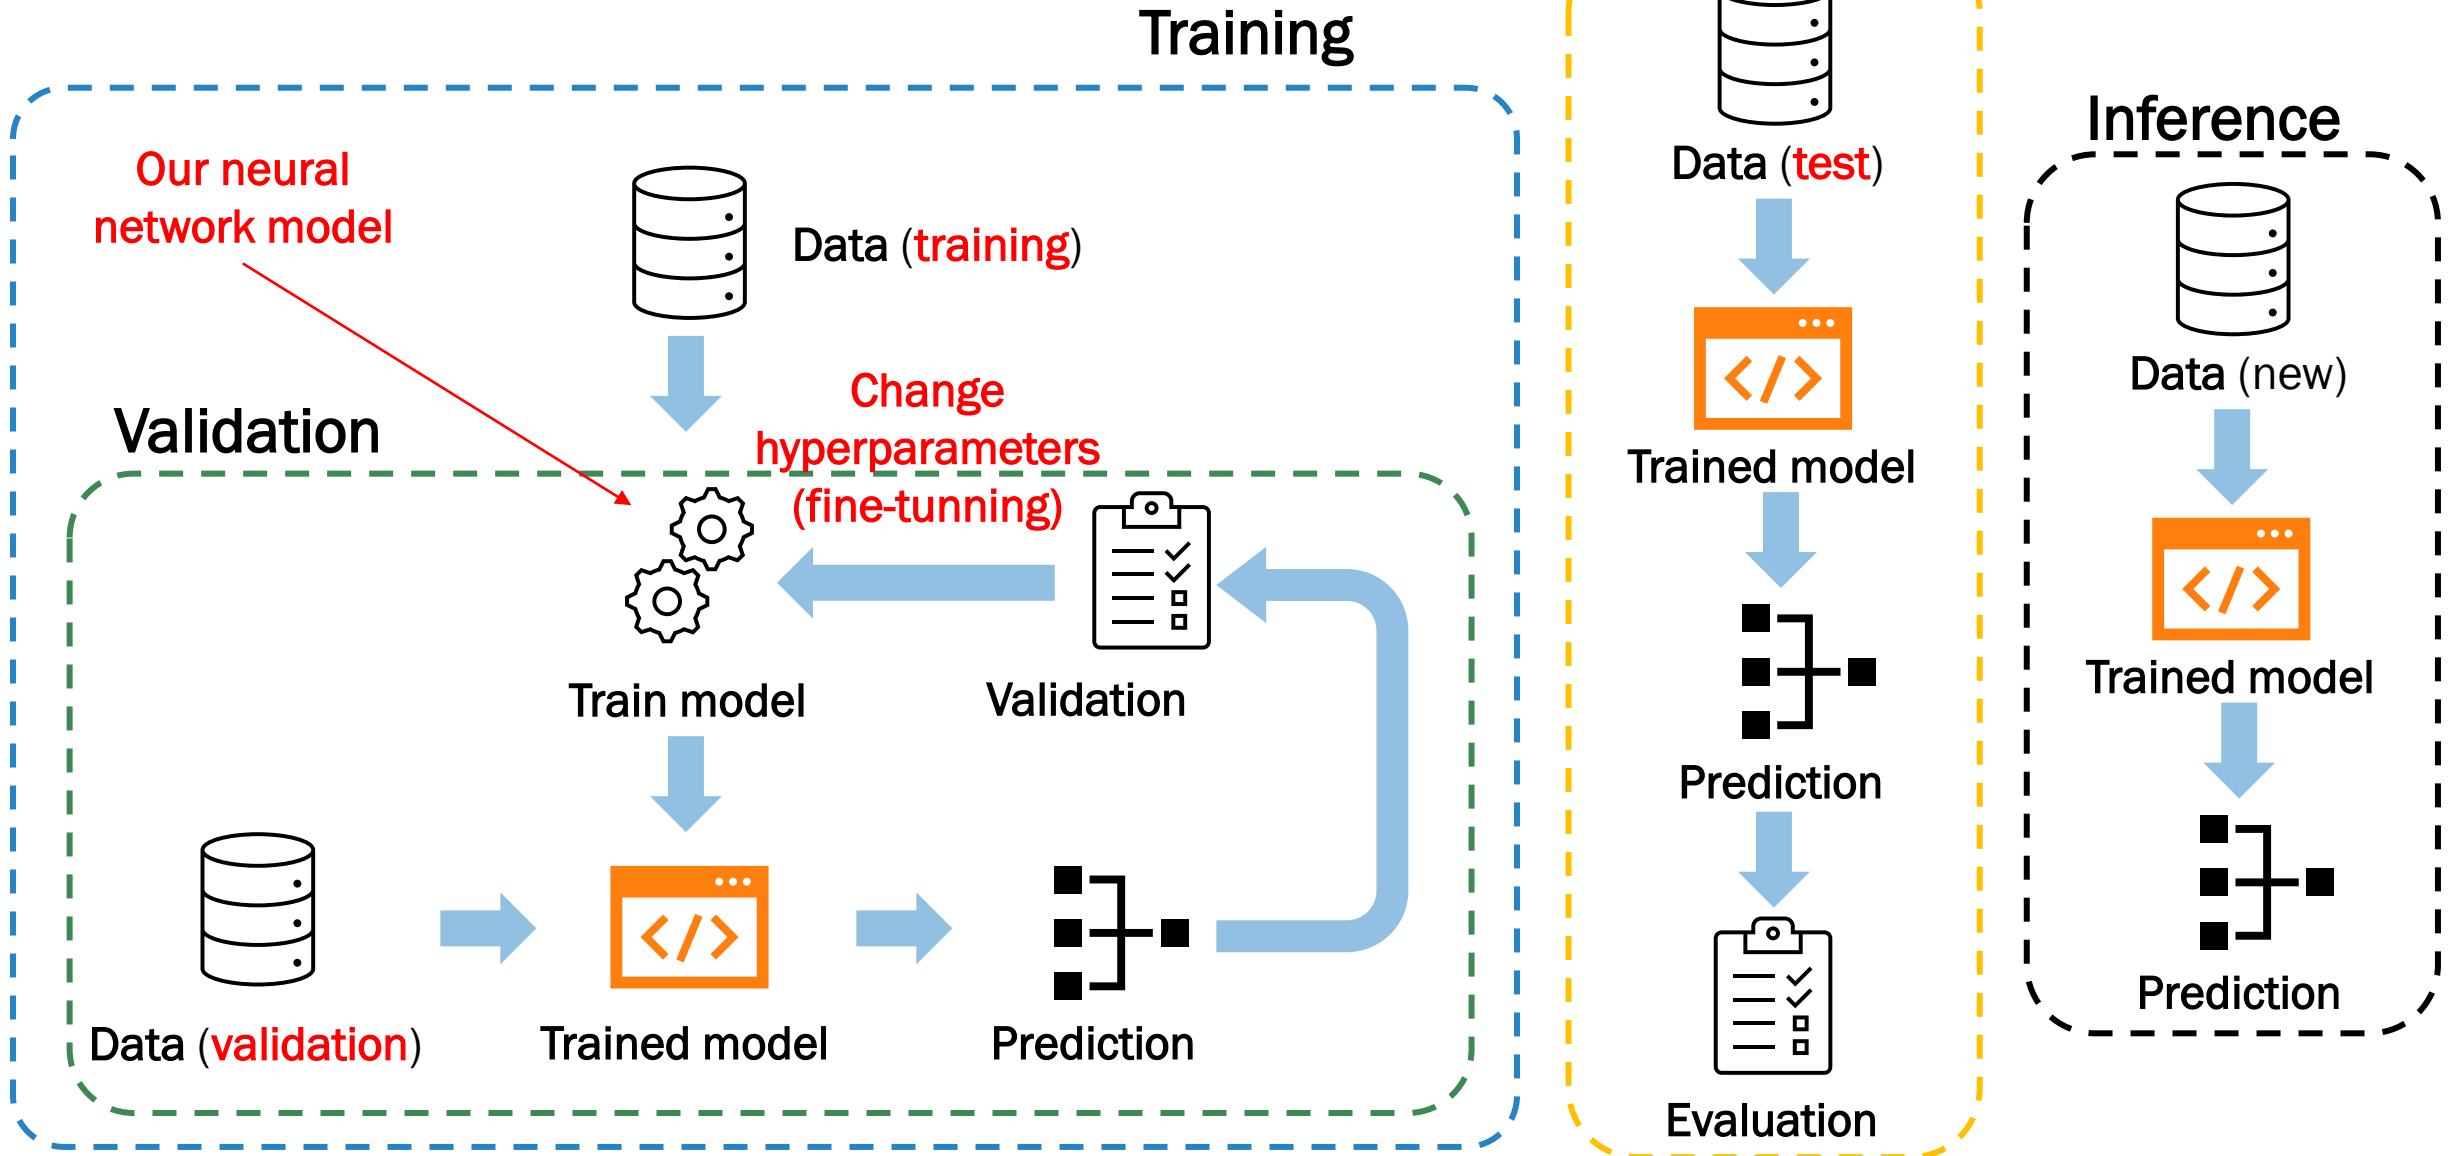


SUPERVISED VS UNSUPERVISED




We have both **supervised** and **unsupervised** neural networks. We will focus the rest of the lecture on **supervised** learning.

MACHINE LEARNING




DATA SPLIT

- Data split



- Know your data:




NEURAL NETWORK EXAMPLE

Input

Input layer

Multiple hidden layers

Output layer



How do we
initialise
weights?

How can we
find the
optimum
weights so that
the error is
minimal?

TRAINING A NEURAL NETWORK

Train model

- Objective: find the optimum weights that will give the lowest **error** [loss]
- Optimisation (Gradient Descent)
 1. Randomly initialise weights
 2. Find how much the weights need to change (by computing the derivatives)
 3. Update weights by taking a small step [learning rate]
 4. Repeat until convergence (until there is no improvement) [number of epochs]

FOR REFERENCE ONLY: LOSS OPTIMISATION

Train model

- Identify (train) a set of weights that will give us the minimum error
- Loss functions: Cross entropy, Mean Squared Error, Mean Absolute Error.
- Needs to be differentiable (the derivative of the function exists for all points).


Binary classification	Multiclass classification	Regression
Loss: Cross Entropy	Loss: Cross Entropy	Loss: Mean Squared Error
Activation: Sigmoid	Activation: Softmax	

BACKPROPAGATION: BACKWARD PROPAGATION OF ERRORS

Train model

- Compute the gradients (computationally expensive!)

How does a small change on the weights affect the final loss $J(W)$?



Watch:

<https://www.youtube.com/watch?v=IN2XmBhILt4&list=PLblh5JKOoLUlxGDQs4LFFD--41Vzf-ME1&index=4>


How much a change in w_2 will affect the total error (loss) $J(W)$?

The partial derivative of the loss with respect to w_2

$$\frac{\partial J(W)}{\partial w_2} = \frac{\partial J(W)}{\partial \hat{y}} * \frac{\partial \hat{y}}{\partial w_2}$$

How much a change in w_1 will affect the total error (loss) $J(W)$?


$$\frac{\partial J(W)}{\partial w_1} = \frac{\partial J(W)}{\partial \hat{y}} * \frac{\partial \hat{y}}{\partial z_1} * \frac{\partial z_1}{\partial w_1}$$




LEARNING RATE

- How much step we take in the gradient?


Train model



Small learning rate



Good learning rate



Large learning rate

FORWARD PROPAGATION AND BACKPROPAGATION



Train model

Input
(flatten)


Input layer

Multiple hidden layers

Output layer



FORWARD PROPAGATION AND BACKPROPAGATION




FORWARD PROPAGATION AND BACKPROPAGATION

Input
(flatten)

Input layer

Multiple hidden layers

Output layer






When poll is active respond at [PollEv.com/esterbonmati](https://pollev.com/esterbonmati)

pollev.com/esterbonmati



What will the NN predict?



Will predict cow

0%


SEE MORE



dict cat or dog


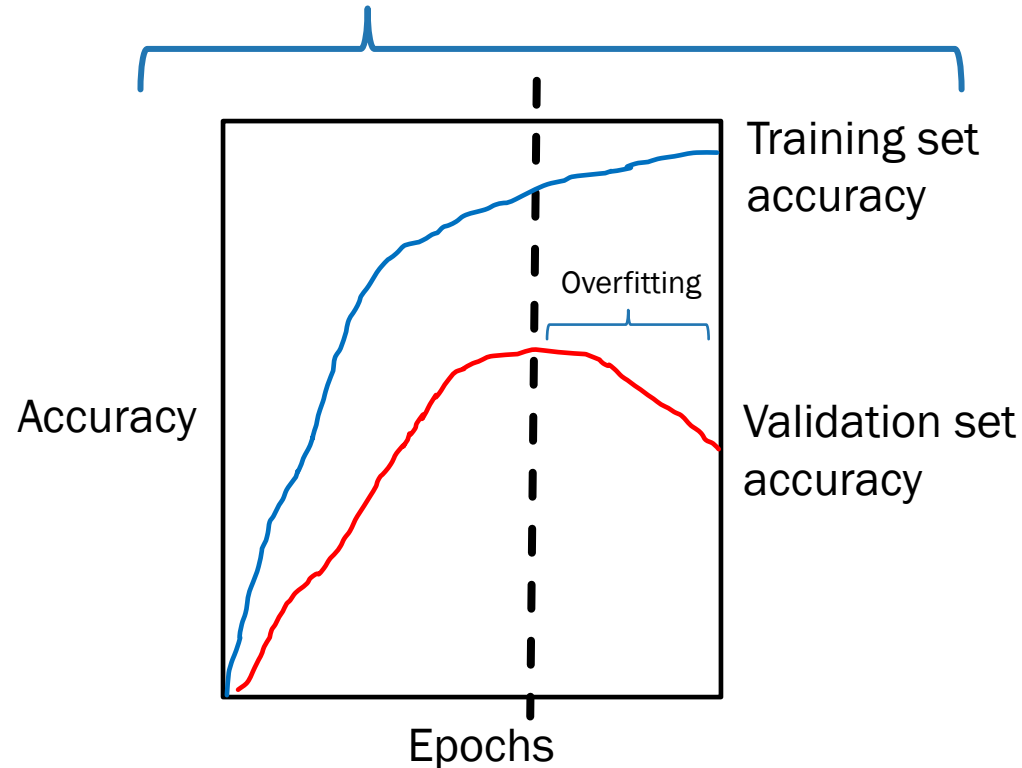
TRAINING

Train model




PROBLEMS: OVERFITTING

- Regularisation
 - **Dropout:** During training, we set some activation to 0 (usually 50%)
 - **Early stopping:** stop training before it overfits the data



TRAINING

Accuracy



Epochs



When poll is active
respond at

[PollEv.com](#)
[/esterbonmati](#)



What would you do next?

Decrease the learning rate

0%

Increase the learning rate

0%

TENSORFLOW & KERAS

- <https://www.tensorflow.org/tutorials/keras/classification>

The screenshot displays the TensorFlow website's 'Basic classification: Classify images of clothing' tutorial page. The left sidebar menu is highlighted with a red dashed box, showing the 'BEGINNER' section with 'Basic image classification' selected. The main content area shows the tutorial title, a table of contents, and buttons to run in Google Colab, view source on GitHub, and download the notebook.

TensorFlow Core

Overview Tutorials Guide Migrate to TF2 TF 1.7

Filter

TensorFlow tutorials

Quickstart for beginners

Quickstart for experts

BEGINNER

ML basics with Keras

Basic image classification

Basic text classification

Text classification with TF Hub

Regression

Overfit and underfit

Save and load

Tune hyperparameters with the Keras Tuner

More examples on keras.io

Load and preprocess data

ADVANCED

Customization

Distributed training

Images

TensorFlow > Learn > TensorFlow Core > Tutorials

Was this helpful?

Basic classification: Classify images of clothing

On this page

- Import the Fashion MNIST dataset
- Explore the data
- Preprocess the data
- Build the model
 - Set up the layers
 - Compile the model
- Train the model
 - Feed the model
- ...

Run in Google Colab

View source on GitHub

Download notebook

This guide trains a neural network model to classify images of clothing, like sneakers and shirts. It's okay if you don't understand all the details; this is a fast-paced overview of a complete TensorFlow program with the details explained as you go.

This guide uses [tf.keras](#), a high-level API to build and train models in TensorFlow.

PYTORCH

- https://pytorch.org/tutorials/beginner/basics/buildmodel_tutorial.html



Table of Contents



Run in Microsoft Learn



Run in Google Colab



Download Notebook



View on GitHub

[Learn the Basics](#) || [Quickstart](#) || [Tensors](#) || [Datasets & DataLoaders](#) || [Transforms](#) || **[Build Model](#)** || [Autograd](#) || [Optimization](#) || [Save & Load Model](#)

BUILD THE NEURAL NETWORK


Neural networks comprise of layers/modules that perform operations on data. The `torch.nn` namespace provides all the building blocks you need to build your own neural network. Every module in PyTorch subclasses the `nn.Module`. A neural network is a module itself that consists of other modules (layers). This nested structure allows for building and managing complex architectures easily.

In the following sections, we'll build a neural network to classify images in the FashionMNIST dataset.

```
import os
import torch
from torch import nn
from torch.utils.data import DataLoader
from torchvision import datasets, transforms
```

ALSO IN MATLAB

- <https://uk.mathworks.com/help/deeplearning/ug/classification-with-a-2-input-perceptron.html>

ProductsSolutionsAcademiaSupportCommunityEventsGet MATLAB

Help CenterSearch Help CenterHelp Center

CONTENTS« Documentation Home« AI, Data Science, and Statistics« Deep Learning Toolbox« Function Approximation, Clustering, and Control« Function Approximation and Clustering« Define Shallow Neural Network ArchitecturesClassification with a Two-Input Perceptron

DocumentationMoreVideosAnswersTrial SoftwareProduct Updates

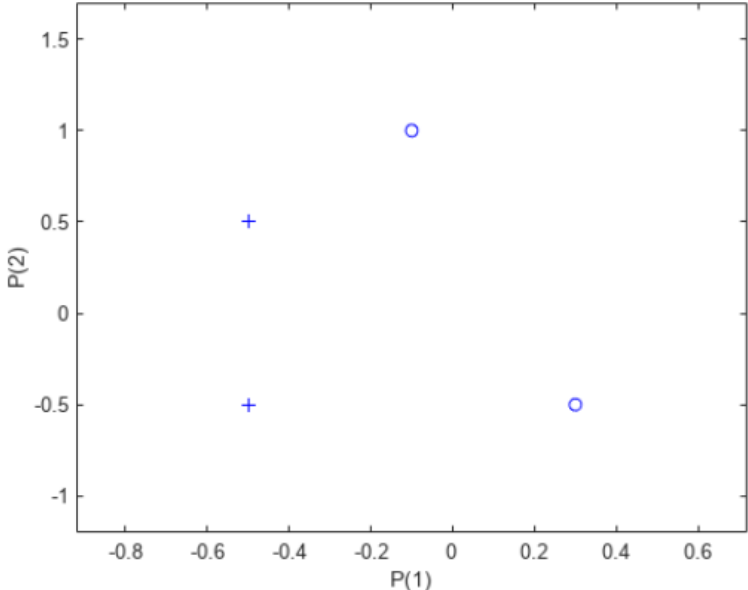
Classification with a Two-Input Perceptron

A two-input hard limit neuron is trained to classify four input vectors into two categories.

Each of the four column vectors in X defines a two-element input vectors and a row vector T defines the vector's target categories. We can plot these vectors with `PLOTPV`.

```
X = [ -0.5 -0.5 +0.3 -0.1; ...
      -0.5 +0.5 -0.5 +1.0];
T = [ 1 1 0 0];
plotpv(X,T);
```

Vectors to be Classified





When poll is active respond at [PollEv.com/esterbonmati](https://pollev.com/esterbonmati)

pollev.com/esterbonmati



After: what is a neural network or artificial neural network?

Nobody has responded yet.

Hang tight! Responses are coming in.

ACKNOWLEDGEMENTS

Acknowledgements: This lecture was inspired by the following material:

- <http://introtodeeplearning.com/>: **MIT Introduction to Deep Learning | 6.S191**
- Neural Networks (Applied AI 21/22 – Dr. Artie Basukoski)